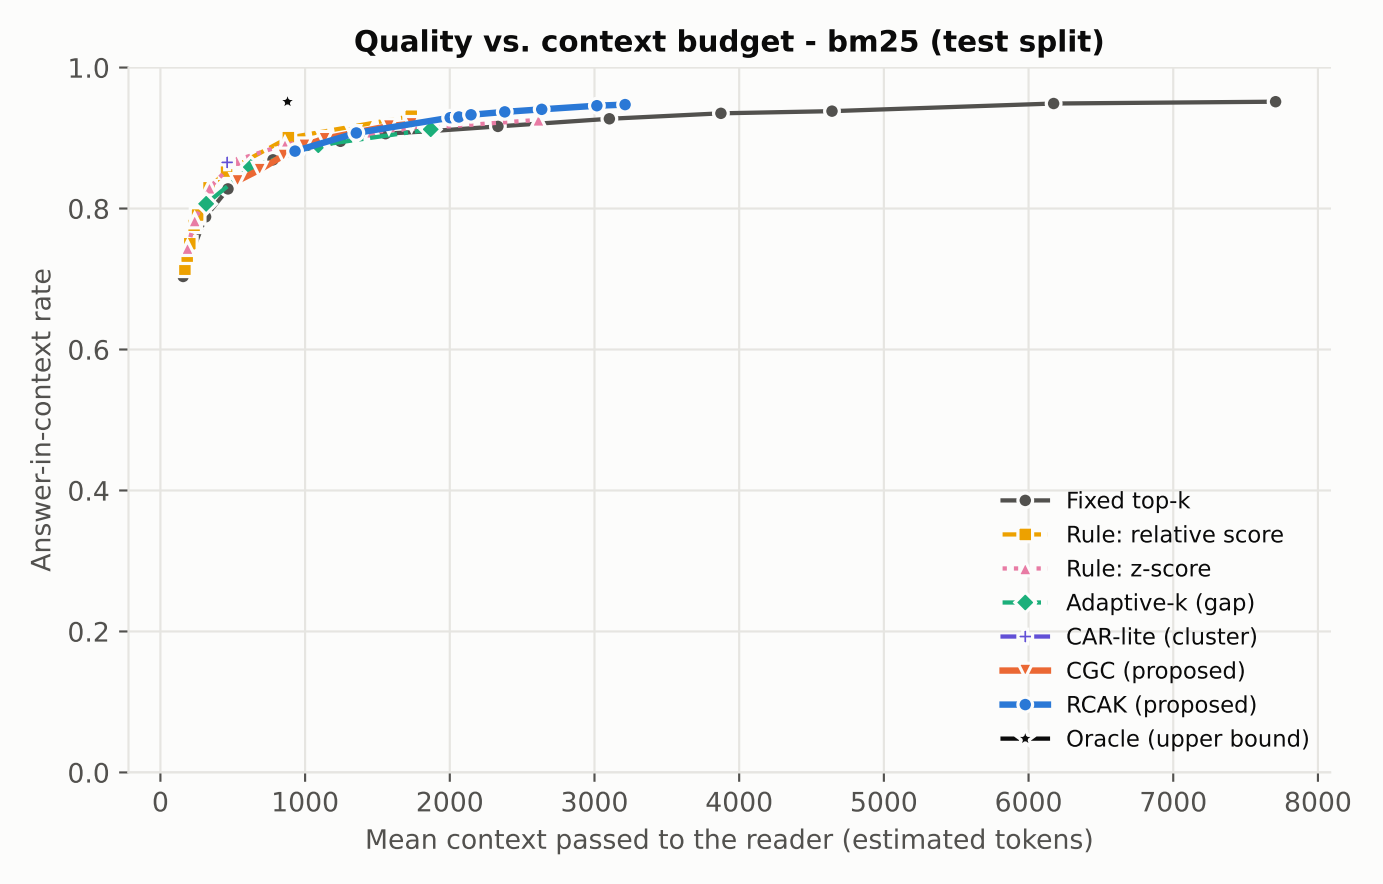
Quality vs. context budget - bm25 (test split)
Answer-in-context rate
1.0
0.8
0.6
0.4
0.2
0.0
0
1000
2000
3000
4000
5000
6000
7000
8000
Mean context passed to the reader (estimated tokens)
Fixed top-k
Rule: relative score
Rule: z-score
Adaptive-k (gap)
CAR-lite (cluster)
CGC (proposed)
RCAK (proposed)
Oracle (upper bound)